

Transformers, Attention and Pre-training


Josh Nguyen
University of Pennsylvania

Notebook: <https://tinyurl.com/4fnu388x>




What is pre-training?

- LLMs are capable of many tasks that can be represented and processed in natural language.
- But before they can do these tasks, they must first be *fluent* in natural language.
- **Pre-training** is the process of making an LLM fluent in human languages.
- This process is called **language modeling**.
(What??)



What is pre-training?



What does it mean to *model* language?

- In the beginning the universe was created.
- Happy families are all alike; every unhappy family is unhappy in its own way.
- Intelligent life on a planet comes of age when it first works out the reason for its own existence.
- Colorless green ideas sleep furiously.
- Contextual nuances highlights the recursive tension inherent to the framework.
- True value lies in embracing the multifaceted tapestry of perspectives.

To **model** language is to have a way of telling how **plausible** a sentence would arise from an actual human language.

We all have an (informal) model of language in our heads.

Probability

- Given a sentence s = “In the beginning the universe was created.”
- What is its probability, $P(s)$?
- Ideally:
 - A human-generated sentence would have high $P(s)$.
 - Other sentences would have much lower $P(s)$.
- So, how we we find this probability?

Sequence of tokens

- A sentence (or a paragraph, a passage, etc.) is represented as a **sequence**.
- Each element in a sequence is a **token**.
- The process of breaking down a sentence is called **tokenization**.

So s = “In the beginning the universe was created.” would be tokenized into

In the beginning the universe was created .

which is a sequence of 8 tokens.

Causal language modeling

Generative process: one token at a time, from left to right.

$$\begin{aligned} &P(\text{In the beginning the universe was created} \text{ . }) \\ &= P(\text{In}) \\ &\times P(\text{the} \mid \text{In}) \\ &\times P(\text{beginning} \mid \text{In the}) \\ &\dots \\ &\times P(\text{.} \mid \text{In the beginning the universe was created}) \end{aligned}$$

Causal language modeling

- So, for a sequence $s = (X_1, X_2, \dots, X_n)$
- The joint probability of this sequence is

$$P(s) = \prod_{i=1}^n P(X_i \mid X_{1:i-1})$$

- If $i = 1$, then $X_{1:i-1} = \emptyset$, so $P(X_i \mid X_{1:i-1}) = P(X_1)$
- So, one marginal probability, and $n - 1$ conditional probabilities.

Causal language modeling

Our language model P can

- Estimate the probability of any sequence. The closer it is to human language, the higher its probability.
- **Also generate:** Given any context s , $P(\cdot | s)$ estimates the probability of *any* possible next token. We can thus sample from $P(\cdot | s)$ “indefinitely.”

Most modern LLMs are causal sequence models.

Generating text from a causal language model

arXiv: 2306.15719

Original abstract

The Magellanic Stream (MS) - an enormous ribbon of gas spanning 140° of the southern sky trailing the Magellanic Clouds - has been exquisitely mapped in the five decades since its discovery. However, despite concerted efforts, no stellar counterpart to the MS has been conclusively identified. This stellar stream would reveal the distance and 6D kinematics of the MS, constraining its formation and the past orbital history of the Clouds. We have been conducting a spectroscopic survey of the most distant and luminous red giant stars in the Galactic outskirts. From this dataset, we have discovered a prominent population of 13 stars matching the extreme angular momentum of the Clouds, spanning up to 100° along the MS at distances of 60–120 kpc. Furthermore, these kinematically-selected stars lie along a $[\alpha/\text{Fe}]$ -deficient track in chemical space from $-2.5 < [\text{Fe}/\text{H}] < -0.5$, consistent with their formation in the Clouds themselves. We identify these stars as high-confidence members of the Magellanic Stellar Stream. Half of these stars are metal-rich and closely follow the gaseous MS, whereas the other half are more scattered and metal-poor. We argue that the metal-rich stream is the recently-formed tidal counterpart to the MS, and speculate that the metal-poor population was thrown out of the SMC outskirts during an earlier interaction between the Clouds. The Magellanic Stellar Stream provides a strong set of constraints - distances, 6D kinematics, and birth locations - that will guide future simulations towards unveiling the detailed history of the Clouds.

Completed by GPT-4

The Magellanic Stream (MS) - an enormous ribbon of gas spanning 140° of the southern sky trailing the Magellanic Clouds - has been exquisitely mapped in the five decades since its discovery. However, despite concerted efforts, no stellar counterpart to the MS has been conclusively identified. This stellar stream would reveal the distance and 6D kinematics of the MS, constraining its formation and the past orbital history of the Clouds. We propose a novel approach to search for the elusive stellar counterpart of the Magellanic Stream. By employing advanced photometric and spectroscopic techniques, combined with state-of-the-art data mining and machine learning algorithms, we aim to systematically analyze and cross-reference multi-wavelength survey data to identify potential candidate stars belonging to the hypothesized stellar stream. Our methodology involves leveraging the distinctive chemical signatures and proper motions associated with stars originating from the Magellanic Clouds, allowing us to differentiate them from the foreground and background contaminants. If successful, this discovery would not only provide essential insights into the origin and evolution of the Magellanic Stream, but also enhance our understanding of the interaction between the Magellanic Clouds and the Milky Way. Furthermore, the proposed approach establishes a powerful framework for investigating other unresolved galactic structures and tracing their underlying dynamical histories.

Completed by LLaMA-2

The Magellanic Stream (MS) - an enormous ribbon of gas spanning 140° of the southern sky trailing the Magellanic Clouds - has been exquisitely mapped in the five decades since its discovery. However, despite concerted efforts, no stellar counterpart to the MS has been conclusively identified. This stellar stream would reveal the distance and 6D kinematics of the MS, constraining its formation and the past orbital history of the Clouds. We find a stellar stream in the MS, with a 6D velocity of 100 km s^{-1} , and a distance of 100 kpc. The stream is 100 pc wide and 1000 pc long, and is consistent with the MS. The stream is 1000 times more massive than the Magellanic Clouds, and is likely to be a tidal stream from the Large Magellanic Cloud.

Completed by AstroLLaMA

The Magellanic Stream (MS) - an enormous ribbon of gas spanning 140° of the southern sky trailing the Magellanic Clouds - has been exquisitely mapped in the five decades since its discovery. However, despite concerted efforts, no stellar counterpart to the MS has been conclusively identified. This stellar stream would reveal the distance and 6D kinematics of the MS, constraining its formation and the past orbital history of the Clouds. We 1) present a new method to select MS candidate stars from wide field optical imaging, 2) apply this method to the MS, 3) compare the results to existing catalogs, and 4) identify the most promising MS candidate stars. We apply our method to 2547 stars in the Gaia EDR3 catalog that have Gaia proper motions consistent with the MS and use these as a pilot sample to optimize our method. We then apply our method to the entire Gaia EDR3 catalog to identify 46,000 stars that are candidates for the MS, and we use the Gaia-ESO Survey to identify 13,000 of these stars that are likely to be MS candidates. We identify 266 stars with (G-K) colours and $[\text{Fe}/\text{H}]$ that are consistent with previous estimates of the MS. These stars have proper motions and radial velocities that are consistent with the MS, and we place them on a 3D map of the MS. We find a previously unknown stellar stream, which we call the Southeastern Stream (SES), in the direction of the Clouds and MS. The SES has an M31 distance of 32.2 ± 0.5 kpc and a 6D velocity of 136 ± 10 km/s. The SES is about 20 times brighter in stars than the MS and about 40 times brighter in stars than the SMC. The SES is also 2.7 times more metal-rich than the MS, which suggests that it is a separate population.

Causal language modeling

How do we find $P(\text{In})$?

- Go out there and **find** all sentences in English.
- **Count** the number of times a sentence starts with In .
- **Normalize** this count to get an estimate of $P(\text{In})$.

Causal language modeling

How do we find $P(\text{the} \mid \text{In})$

- Go out there and **find** all sentences in English.
- **Count** the number of times a sentence starts with `In` .
- **Count** the number of times such a sentence is followed by `the` .
- **Normalize** to get the conditional probability.

Problem!!

- Let's say the English language has about 100,000 words.
- For a sequence of **2** words
 - We need to store every combination of X_1 and X_2 .
 - So we need to store **100,000²** probabilities
- For a sequence of **n** words
 - We need to store **100,000ⁿ** probabilities!!

Solution: Train a statistical model such much fewer parameters.

Evaluating a language model

- Any causal language model has the form $P(\cdot | \cdot)$.
- How do we evaluate competing models of our data? Likelihood.
- For a sequence s , $P(s)$ is the likelihood of our model.
- Thus “better” models would give higher $P(s)$ for each of our sentences.

Evaluating a language model

- Recall that

$$P(s) = \prod_{i=1}^n P(X_i \mid X_{1:i-1}) = P(X_1) \times P(X_2 \mid X_1) \times \dots \times P(X_n \mid X_{1:n-1}).$$

- **Problem!!!** Longer sequences tend to have much lower likelihoods.

Accounting for sequence length

- We need to normalize likelihood by sequence length.
- First, turn likelihood into log-likelihood

$$\log P(s) = \sum_{i=1}^n \log P(X_i \mid X_{1:i-1})$$

- Second, divide by the number of tokens

$$\frac{1}{n} \log P(s) = \frac{1}{n} \sum_{i=1}^n \log P(X_i \mid X_{1:i-1})$$

- Third, exponentiate its negative, and we have the **perplexity score**

$$\text{PPL}(s) = \exp \left(-\frac{1}{n} \log P(s) \right).$$

Exercise time!

Number of tokens used for calculation: 7

Actual likelihood of each token

```
logP( the | In) = -1.8271
logP( beginning | In the) = -4.8863
logP( the | In the beginning) = -3.5942
logP( universe | In the beginning the) = -6.0115
logP( was | In the beginning the universe) = -0.7214
logP( created | In the beginning the universe was) = -0.5190
logP(. | In the beginning the universe was created) = -0.4145
```

Log likelihood: -17.9740

Perplexity: 13.0360

Notebook: <https://tinyurl.com/4fnu388x>



Training a language model

$$P(\text{next_token} \mid \text{some context})$$

- Estimating the conditional $P(\cdot \mid \cdot)$ is doable; we just need to count.
- But it is *not* feasible due to the complexity of the space.
- Instead, we will train a model to approximate $P(\cdot \mid \cdot)$.
- What kinds of model can do this well?

N-gram language model

- Markov assumption: ignore the entire context but the latest (N-1) tokens.

$$P(\text{beginning} \mid \text{In the}) = P(\text{beginning} \mid \text{the})$$

- With 2-grams, complexity is now just **100,000²** probabilities.
- Easy to understand and calculate, fully interpretable.
- Longer context tends to fit better. But overall does not generalize very well.


	Unigram	Bigram	Trigram
Perplexity	962	170	109

Jurafsky & Martin 2026, *Speech and Language Processing*, 3rd ed., p. 10.


N-gram models trained on 38M words from WSJ

Neural network language model

- Train a model P that approximates prob. of the next token, $P(\cdot | \cdot)$.
- Find model parameters θ that minimize the perplexity of the training set.



Predicting the next token



Embedding a token

- The embedding function f
 - Takes any token in our vocabulary, like In
 - Then transforms it into a dense vector $f(\text{In})$
- The crux of our language model is to design a good embedding function.
- Several decisions
 - Number of dimensions: 512, 768, 1024, 2048, etc.
 - How it is composed and learned: feedforward, convolution, recurrent, attention, etc.

Attention: choosing the right context

Imagine generating this sentence


The trophy can't fit inside the trunk because it

How should we determine the embedding of **it** ?

- **it** is related to certain words than others.
- Namely, **trophy** and **trunk**, depending on the content.
- Either way, **it** should mostly “attend” to **trophy** and **trunk**.

Attention: choosing the right context

If we were to display it like a graph, it would be something like this



So, if the embedding of `it` may rely on the context before it, the parts of that context that matter the most are `trophy` and `trunk`.

Attention: choosing the right context

Computationally, the attention “head” has three parametrized components


- **Query**: a vector that represents `it`.
- **Key**: a set of vectors that represent the context, including `trophy` and `trunk`.
- **Value**: a set of vectors for all tokens to compute the output.

The final embedding of `it` is the weighted average of all **values**, the weights being determined by the similarity between the **query** and **keys**.

$$A = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

Transformer block


- A transformer block includes **attention** and other other components, such as **residual** stream, layer **norm** and **feedforward**.
- Each token “attends” to everything else before it, just with different weights.
- Extremely useful for sequence modeling.



Vaswani *et al.* 2017, Attention Is All You Need, arXiv: 1706.03762

Pre-training in practice

- Find model parameters that minimize perplexity on a training set.
- Huge engineering and computational endeavor. Done mostly by big players.
- Lots of data quality control needed.



Olmo Team *et al.* 2026,
Olmo 3, arXiv: 2512.13961

Continual pre-training

- Pre-training from scratch is most likely beyond our means.
- But **continual pre-training** can be.
- Idea: take a pre-trained LLM, continue training it on our dataset.
- Goal: make the model more **fluent** in our domain than it already is.


AstroLLaMA : Towards Specialized Foundation Models in Astronomy

Tuan Dung Nguyen^{1, 2*}, Yuan-Sen Ting^{2, 3*}, Ioana Ciucă^{2*}
Charles O'Neill^{2†}, Ze-Chang Sun^{4†}, Maja Jabłońska^{2†}, Sandor Kruk^{5†}
Ernest Perkowski⁵, Jack Miller², Jason Li⁶, Josh Peek⁷, Kartheik Iyer⁸,
Tomasz Rózański^{2, 9}, Pranav Khetarpal¹⁰, Sharaf Zaman², David Brodrick²
Sergio J. Rodríguez Méndez², Thang Bui², Alyssa Goodman¹¹, Alberto Accomazzi¹²,
Jill Naiman¹³, Jesse Cranney², Kevin Schawinski¹⁴, UniverseTBD

arXiv: 2309.06126

Continual pre-training

- 326,238 abstracts from arXiv astro-ph (1992-04 to 2023-07).
- 95M tokens, 77M used for training.
- Pre-trained model: LLaMA-2, 6.7B parameters.
- Model fit evaluated by perplexity.



“Extrinsic” evaluation

- Perplexity measures model fit to our data, regardless of the task. Thus it is called an **intrinsic metric**.
- But models must also be evaluated **extrinsically**, i.e., based on tasks that we actually use them for.
- What kinds of task can we do with a pre-trained LLM?
 - Generating text
 - Embedding text for similarity and retrieval

Generating more plausible text

arXiv: 2306.15719

Original abstract

The Magellanic Stream (MS) - an enormous ribbon of gas spanning 140° of the southern sky trailing the Magellanic Clouds - has been exquisitely mapped in the five decades since its discovery. However, despite concerted efforts, no stellar counterpart to the MS has been conclusively identified. This stellar stream would reveal the distance and 6D kinematics of the MS, constraining its formation and the past orbital history of the Clouds. We have been conducting a spectroscopic survey of the most distant and luminous red giant stars in the Galactic outskirts. From this dataset, we have discovered a prominent population of 13 stars matching the extreme angular momentum of the Clouds, spanning up to 100° along the MS at distances of 60–120 kpc. Furthermore, these kinematically-selected stars lie along a $[\alpha/\text{Fe}]$ -deficient track in chemical space from $-2.5 < [\text{Fe}/\text{H}] < -0.5$, consistent with their formation in the Clouds themselves. We identify these stars as high-confidence members of the Magellanic Stellar Stream. Half of these stars are metal-rich and closely follow the gaseous MS, whereas the other half are more scattered and metal-poor. We argue that the metal-rich stream is the recently-formed tidal counterpart to the MS, and speculate that the metal-poor population was thrown out of the SMC outskirts during an earlier interaction between the Clouds. The Magellanic Stellar Stream provides a strong set of constraints - distances, 6D kinematics, and birth locations - that will guide future simulations towards unveiling the detailed history of the Clouds.

Completed by LLaMA-2

The Magellanic Stream (MS) - an enormous ribbon of gas spanning 140° of the southern sky trailing the Magellanic Clouds - has been exquisitely mapped in the five decades since its discovery. However, despite concerted efforts, no stellar counterpart to the MS has been conclusively identified. This stellar stream would reveal the distance and 6D kinematics of the MS, constraining its formation and the past orbital history of the Clouds. We find a stellar stream in the MS, with a 6D velocity of 100 km s⁻¹, and a distance of 100 kpc. The stream is 100 pc wide and 1000 pc long, and is consistent with the MS. The stream is 1000 times more massive than the Magellanic Clouds, and is likely to be a tidal stream from the Large Magellanic Cloud.

Completed by GPT-4


The Magellanic Stream (MS) - an enormous ribbon of gas spanning 140° of the southern sky trailing the Magellanic Clouds - has been exquisitely mapped in the five decades since its discovery. However, despite concerted efforts, no stellar counterpart to the MS has been conclusively identified. This stellar stream would reveal the distance and 6D kinematics of the MS, constraining its formation and the past orbital history of the Clouds. We propose a novel approach to search for the elusive stellar counterpart of the Magellanic Stream. By employing advanced photometric and spectroscopic techniques, combined with state-of-the-art data mining and machine learning algorithms, we aim to systematically analyze and cross-reference multi-wavelength survey data to identify potential candidate stars belonging to the hypothesized stellar stream. Our methodology involves leveraging the distinctive chemical signatures and proper motions associated with stars originating from the Magellanic Clouds, allowing us to differentiate them from the foreground and background contaminants. If successful, this discovery would not only provide essential insights into the origin and evolution of the Magellanic Stream, but also enhance our understanding of the interaction between the Magellanic Clouds and the Milky Way. Furthermore, the proposed approach establishes a powerful framework for investigating other unresolved galactic structures and tracing their underlying dynamical histories.

Completed by AstroLLaMA

The Magellanic Stream (MS) - an enormous ribbon of gas spanning 140° of the southern sky trailing the Magellanic Clouds - has been exquisitely mapped in the five decades since its discovery. However, despite concerted efforts, no stellar counterpart to the MS has been conclusively identified. This stellar stream would reveal the distance and 6D kinematics of the MS, constraining its formation and the past orbital history of the Clouds. We 1) present a new method to select MS candidate stars from wide field optical imaging, 2) apply this method to the MS, 3) compare the results to existing catalogs, and 4) identify the most promising MS candidate stars. We apply our method to 2547 stars in the Gaia EDR3 catalog that have Gaia proper motions consistent with the MS and use these as a pilot sample to optimize our method. We then apply our method to the entire Gaia EDR3 catalog to identify 46,000 stars that are candidates for the MS, and we use the Gaia-ESO Survey to identify 13,000 of these stars that are likely to be MS candidates. We identify 266 stars with (G-K) colours and $[\text{Fe}/\text{H}]$ that are consistent with previous estimates of the MS. These stars have proper motions and radial velocities that are consistent with the MS, and we place them on a 3D map of the MS. We find a previously unknown stellar stream, which we call the Southeastern Stream (SES), in the direction of the Clouds and MS. The SES has an M31 distance of 32.2 ± 0.5 kpc and a 6D velocity of 136 ± 10 km/s. The SES is about 20 times brighter in stars than the MS and about 40 times brighter in stars than the SMC. The SES is also 2.7 times more metal-rich than the MS, which suggests that it is a separate population.

Embedding texts for retrieval

- Used 2 LLMs to embed **10,000 abstracts** using token embedding.
- For each pair of abstracts, calculated their cosine similarity.
- **GPT** thinks most abstracts are **very similar** (domain generality).
- **AstroLLaMA** thinks there's **much more variance** in similarity (suggesting domain specialization).



Paper 1: Astrophysical gyrokinetics: kinetic and fluid turbulent cascades in magnetized weakly collisional plasma

Paper 2: Comment on modified Coulomb law in a strongly magnetised vacuum

GPT-3 cosine similarity score: **78.5%**

AstroLLaMa cosine similarity score: 36.3%


Paper 1: A Spitzer census of the IC 348 nebula

Paper 2: Sequential and spontaneous star formation around the mid-infrared halo HII region KR 14

GPT-3 cosine similarity score: 82.4%

AstroLLaMa cosine similarity score: **92.8%**

What's next?



AstroLLaMA-Chat: Scaling AstroLLaMA with Conversational and Diverse Datasets

ERNEST PERKOWSKI*,¹ RUI PAN*,² TUAN DUNG NGUYEN,³ YUAN-SEN TING,^{4,5,6,7} SANDOR KRUK,¹ TONG ZHANG,⁸
CHARLIE O'NEILL,⁹ MAJA JABLONSKA,⁴ ZECHANG SUN,¹⁰ MICHAEL J. SMITH,¹¹ HUILING LIU,¹² KEVIN SCHAWINSKI,¹³
KARTHEIK IYER,¹⁴ IOANA CIUCĂ,^{4,5} AND UNIVERSETBD

Supervised fine-tuning

What is the evidence for cosmological parity violation?

AstroLLaMA-Chat

The study provides indirect evidence for cosmological parity violation using observables like cosmic microwave background radiation and galaxy clustering statistics. Specifically, they analyzed the correlation between CMB quadrupole amplitude and dipolar power, finding a weak preference for even multipole moments. They also detected a slightly stronger signal when comparing odd versus even multipoles. Both results suggest deviations from statistical isotropy, potentially indicating a deviation from parity conservation in the early Universe. Further analysis will help determine if these anomalous signals are statistically significant and robust against systematic effects. If confirmed, this could revolutionize our understanding of the origins and evolution of the Cosmos.

GPT-4

As of my last update in April 2023, cosmological parity violation is a theoretical concept that suggests a fundamental asymmetry in the laws of physics, particularly in how they apply to mirrored configurations of physical systems. In the context of cosmology, this would imply that the universe might not be identical to its mirror image, violating what is known as "parity symmetry". However, it's important to note that this is a highly speculative area and evidence for such a violation at the cosmological scale, if it exists, is subtle and subject to ongoing research and debate. Here are some key points regarding the potential evidence and considerations: 1. Cosmic Microwave Background (CMB) Anomalies [abridged], 2. Baryon asymmetry [abridged], 3. Polarization of light from distant galaxies [abridged], 4. Gravitational lensing [abridged], 5. Gravitational waves [abridged], 6. Limits from observational data [abridged], 7. Role of advanced observations and simulations [abridged]

LLaMA-2-7b-Chat

Cosmological parity violation refers to the idea that the laws of physics may be different in different directions in space, or that the symmetries of the universe may be broken in a way that depends on the direction in which we look. While this idea may seem strange, there is some evidence that it could be true. Here are some of the key lines of evidence: 1. CMB temperature anisotropies [abridged], 2. Polarization of the CMB [abridged], 3. Large-scale structure of the universe [abridged], 4. Gravitational lensing [abridged], 5. Nuclear scattering [abridged], 6. Neutrino experiment [abridged], 7. High-energy cosmic rays [abridged], 8. Gamma-ray bursts [abridged]